




# Memory and Memory Caches

Joe Sventek

# Commodity Hardware Today

- All CPUs connected via a common bus (the Front Side Bus) to the Northbridge
- Northbridge contains the memory controller
- To reach other system devices, Northbridge must communicate with Southbridge
- Southbridge (aka I/O bridge) provides access to a variety of different buses




# Consequences of this structure

- All interCPU communication must travel over the same bus used to communicate with Northbridge
- All communication with RAM must pass through Northbridge
- The RAM has only a single port
- Communications between a CPU and a device attached to Southbridge is routed through Northbridge

# Bottlenecks

- DMA access creates contention for the bandwidth of Northbridge, as DMA requests compete with RAM access from the CPUs
- The single port to the memory (which may be one or more channels) is another source of contention; with limited bandwidth available, Northbridge must schedule memory access in ways that minimize delays

# External controllers



- On more expensive systems, Northbridge does not contain the memory controller; instead, it can be connected to a number of external memory controllers
- This provides enhanced memory bus bandwidth
- System still limited by the internal bandwidth of the Northbridge



## Static RAM (SRAM)

- Maintaining the state of a memory cell requires constant power
- The memory cell state is available for reading almost immediately once the word access line is raised
- The memory cell state is stable, no refresh cycles are needed.

## Dynamic RAM (DRAM)

- Due to leakage, each cell must be constantly refreshed (~ every 64ms); during refresh, no access to the memory is possible
- The memory cell state is not immediately available for reading
- The chip real estate needed for a DRAM cell is many times smaller than for SRAM; thus can pack many more DRAM cells on a chip, and DRAM is cheaper

# Types of RAM

# Addressing Northbridge contention

- Use SRAM – not economical, SRAM as fast as current CPU cores is orders of magnitude more expensive than DRAM
- Introduce a small amount of high-speed SRAM to complement large amounts of DRAM
- The SRAM could be part of the addressable memory of the machine, forcing the OS to optimally distribute instructions/data to make best use of the SRAM
  - Logistically a nightmare, as requires each process to administer in software the allocation of the SRAM memory region.

# Caches

- Make the SRAM a resource that is transparently used and administered by the processors
- SRAM is used to make temporary copies of (i.e. to cache) data in main memory which is likely to be used soon by the processor
- Possible because program code and data exhibit spatial and temporal locality
  - Spatial – loops in code, sequential access to data
  - Temporal – likely that the same data/instructions will be reused before long



# Example speedup


- Assume access to main memory takes 200 cycles and access to cache memory takes 15 cycles.
- If code uses 100 data elements 100 times:
  - If no cache, then code will spend 2,000,000 cycles on memory operations
  - If all data is in the cache, code will spend 168,500 cycles
  - This is an improvement of 91.5%

# Can the working set fit in the cache?

- Cache size  $\sim 10^{-3} \times$  DRAM size (e.g. 4MB cache for a system with 4 GB of DRAM)
- Working set for all processes is bound to be larger than the cache
- Need a set of good strategies to determine what should be cached at any given time
- Since not all data of the working set is used at exactly the same time, can use techniques to temporarily replace some data in the cache with other data


# Where do caches fit in the architecture?

Minimum




- All loads and stores go through the Cache
- The Bus is the FSB
- Assume Northbridge is there to facilitate CPU-mem communications

More Typical



- Separate instruction and data caches
- Each level is larger, but slower, than the previous level.

# What about multicore?



- Each core has its own level 1 cache
- All cores on a processor share level 2 and level 3 caches
- The different processors do not share any caches

# Cache operations

- All data read or written by a CPU core is stored in the cache
- When a CPU needs a data word, the caches are searched first
- The cache cannot contain the contents of all of main memory, and all memory addresses are cacheable
- Each cache entry is *tagged* using the address of the data word in main memory
- The address can be physical or virtual, depending upon the cache implementation




# Cache operations (cont)

- Since the tag requires space in addition to the actual contents of the data word, it is inefficient to choose a word (e.g. 32-bits) as the granularity of the cache.
- Thus, entries stored in the cache are lines of several contiguous words
- Typical line size today is 64 bytes
- If the memory bus is 64bits wide, this means 8 transfers per cache line.
- This transport mode is supported efficiently by modern memories (DDR)



# Cache operations (cont)



- When memory content is needed by the CPU, the entire cache line is loaded into the L1d
- In the figure above,
  - The cache line size is  $2^O$  bits; the low O bits are the offset into the cache line
  - The next S bits select the Cache Set (see section 3.3)
  - The remaining T bits are used to distinguish all addresses with the same S part of the address
  - The S bits do not have to be stored in the cache since they are the same for all cache lines in the same set

# Cache operations (cont)

- When an instruction modifies memory, the CPU has to load a cache line first because no instruction modifies an entire cache line at once
- It is not possible for a cache to hold partial cache lines
- A cache line that has been written to and which has not been written back to main memory is “dirty”
- Once it is written to main memory, the dirty flag is cleared

# Interactions between caches

- To load new data in a cache, it is almost always first necessary to make room in the cache
- An eviction from L1d pushes the cache line down into L2 (uses the same cache size)
- This means that room must be made in L2, evicting the content into L3.
- Eviction from L3 means that the cache line must be written to main memory
- Exclusive cache – cache line is only in a single cache
- Inclusive cache – each cache line in L1d is also in L2

# Cache management

- CPUs may manage the caches in any way they wish as long as the processor architecture memory model is not violated
- For symmetric multi-processor (SMP) systems, the CPU caches cannot work independently from each other.
  - All processors are supposed to see the same memory content at all times
  - Without collusion, a processor would not see the content of dirty cache lines in other processors

# Cache coherency protocols

- Each CPU must be able to detect when another CPU wants to read or write a certain cache line (snooping)
- If write access is detected and the CPU has a clean copy of the cache line in its cache, it must mark the cache line as invalid; thus, future references by the detecting CPU will require reloading the cache line
- If read access is detected and the CPU has a clean copy, no action is needed

# Cache coherency protocols (cont)


- If read or write access is detected and the CPU has a dirty copy of the cache line in its cache
  - Detecting CPU sends the dirty cache line directly to the requesting CPU
  - Detecting CPU writes the cache line to main memory (or the memory controller is supposed to notice the direct transfer and store the updated cache line in memory)
  - If write access was detected, the detecting CPU then invalidates its local cache line
- See 3.3.4 on the MESI protocol in the handout.



# Relative costs

To Where	Cycles
Register	$\leq 1$
L1d	$\sim 3$
L2	$\sim 14$
Main Memory	$\sim 240$

- Costs provided by Intel for a Pentium M



- L1d is  $2^{13}$  bytes
- L2 is  $2^{20}$  bytes
- No L3 cache